yarnnn

Cloud-based Multi-Agent Framework for Solo-Founders

Pre Seed: 2026

AI Agents that work while you sleep.

yarnnn.com · Kevin Kim · seulkim88@gmail.com

It looks like the big players will own everything.

Claude has coding agents, desktop, cowork.
ChatGPT has memory, Codex, Sora.
The market assumes they'll consume every layer.

But every platform cycle looks that way — until it doesn't.

Google didn't become Salesforce.
Facebook didn't become Shopify.
AWS didn't become Datadog.

The Honest Bet

History never repeats, but it does rhyme. Our thesis is that even with the AI tech revolution, the application layer is yet to come.

2008: Google	→	Looked like it would own everything on the web. It didn't become Salesforce.
2012: Facebook	→	Looked like it would own all of social commerce. It didn't become Shopify.
2015: AWS	→	Looked like it would own every cloud application. It didn't become Datadog.
2026: LLMs	→	Look like they'll own every layer. But code was the easy case. Work context is the hard one.

The application layer of the AI platform shift is yet to come, and we're taking the bold bet that it will be Agent-native, and the work economy will shift from human → agents

Proof of Demand

OpenClaw proved the appetite is real.

17,830

GitHub stars in 24 hours

Fastest single-day growth in GitHub history.

But 95% of people can't use it.

High technical barrier, privacy, hardware setup

> *"The first time I've felt like I'm living in the future since ChatGPT launched."*
>
> — ClawdBot user

- ✗ Requires VPS provisioning (only 5% can do this)
- ✗ 200+ security leaks detected (GitGuardian)
- ✗ No project isolation or recurring workflows

Product Architecture Solution / Moat

Product: Meet yarnnn's Agent Framework

Multi-Agent Framework. Just setup and work

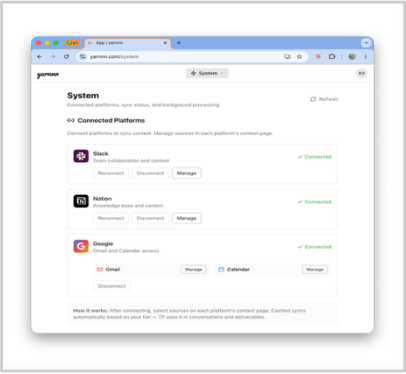

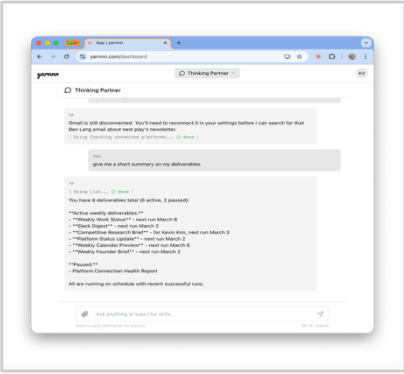

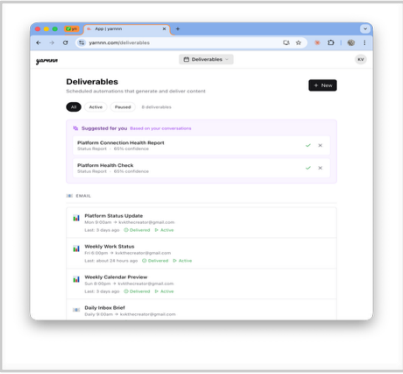

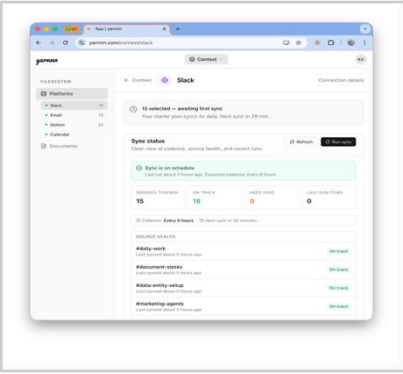

T+0s

Dump, Connect

Tell us what you're working on. System extracts work scopes. No configuration.

T+30s

Agent Setup

Projects auto-created with Agent scaffolding, context loading, work boundary creations

Days

Accumulate

Every sync deepens knowledge. Every edit refines quality. Teams coordinate.

Weeks

Compound

PM coordination matures. Cross-phase handoffs refine. Output quality compounds.

Built for coordinated, autonomous team execution

Every purpose-built layer compounds — creating recursive feedback between agent output and delivery

Perception Pipeline
Slack
Notion

Knowledge Base
Unified content
Retention + provenance

Workspace Filesystem
Charter + memory
Agent workspaces

Agent Teams
8 types, PM dispatch
3-tier pulse

Output Gateway
HTML compose
PDF · PPTX · XLSX

Why this can't be bolted on:

Big Players, LLM Providers

No value chain from description to compounding.
Each session isolated. No team coordination.

Agent Startups

Can execute but can't accumulate.
No Phases 5-7 (compounding).
Every run starts fresh.

Workspace AI

Trapped in one platform. Can't compose cross-platform agent teams.
No recursive perception.

Product

Pick what you want delivered.

13 task types across 5 categories. Each one is a concrete deliverable with a pre-built multi-agent pipeline behind it.

Category	Deliverable	Pipeline (agents)	Default cadence
INTELLIGENCE	Competitive Intel Brief	Research → Content	Weekly
	Market Research Report	Research → Content	Monthly
	Industry Signal Monitor	Marketing → Research	Weekly
OPERATIONS	Meeting Prep Brief	CRM → Research	On-demand
	Stakeholder Update	Research → Content	Monthly
	Project Status Report	Slack → CRM → Content	Weekly
PLATFORM	Slack Recap	Slack Bot	Daily
	Notion Sync Report	Notion Bot	Weekly
CONTENT	Content Brief / Blog Draft	Research → Content	On-demand
	Launch Material	Marketing → Content	On-demand
TRACKING	GTM Tracker	Marketing → Content	Weekly

Users pick deliverables, not agents. Each task type has a pre-meditated pipeline — deterministic, predictable cost, no runtime orchestration overhead. Custom tasks use improvised TP orchestration.

Product

8 Agent Types. Each one hired for a job.

Deterministic (close-to) agent configurations, with individual memory, self-assessment, and collaboration

Briefer

Recurring digests from platform activity. Your daily pulse on what happened.

Monitor

Watches for changes and patterns. Alerts on what matters, ignores what doesn't.

Researcher

Deep dives with web search. Finds context beyond your connected platforms.

Analyst

Cross-references sources. Surfaces patterns humans miss across platforms.

Drafter

Produces first drafts from briefs. Turns PM instructions into starting points.

Writer

Polished final output. Blog posts, reports, memos that ship as-is.

Planner

Roadmaps, strategies, structured plans. Turns goals into actionable sequences.

Scout

Tracks competitors and market signals. Watches the outside while you work inside.

Project Manager Agent

Coordinates contributor phases, quality-gates assembly, steers the team.
One per project.

Product

From Context to Autonomous Delivery

Each phase feeds the next. The system compounds — competitors starting fresh can't catch up.

Context

Connect Slack + Notion. Content syncs and accumulates over time.

Inference

System reads your work. Extracts identity, brand, scopes from docs + chat.

Task Type Selection

Pick deliverables from the catalog. Each has a pre-built pipeline.

Pipeline Execution

Multi-agent pipeline runs mechanically. Each step hands off to the next.

Autonomous Delivery

Output compounds weekly.
Feedback refines agents.
Quality improves with use.

What makes our architecture unique:

Pre-meditated pipelines

Each task type carries a defined execution plan. Agents run in sequence with explicit handoffs. Deterministic cost (~$0.06–$0.16/run). No runtime AI orchestration needed.

Filesystem IS the architecture

TASK.md + AGENT.md define everything. No hidden state. Agents read the same files humans can edit. Workspace accumulates — every run adds to the knowledge substrate.

Knowledge compounds, not resets

Every sync deepens context. Every edit trains preferences. 90 days of accumulated intelligence can't be replicated by starting over. Switching cost is structural.

Solution: Moat

Team Intelligence Compounds. Switching Costs Follow.

Every agent run makes the next one better. Every edit teaches the team. Every sync deepens knowledge.

Three moat mechanics:

1. **Accumulation:** Every iteration cycle deepens cross-platform understanding. Agent teams don't start from zero.

2. **Coordination:** PM learns which phases work, which handoffs succeed. Team intelligence compounds independently of individual agent improvement.

3. **Irreversibility:** 90 days of accumulated team intelligence — agent memory, PM patterns, learned preferences — can't be replicated

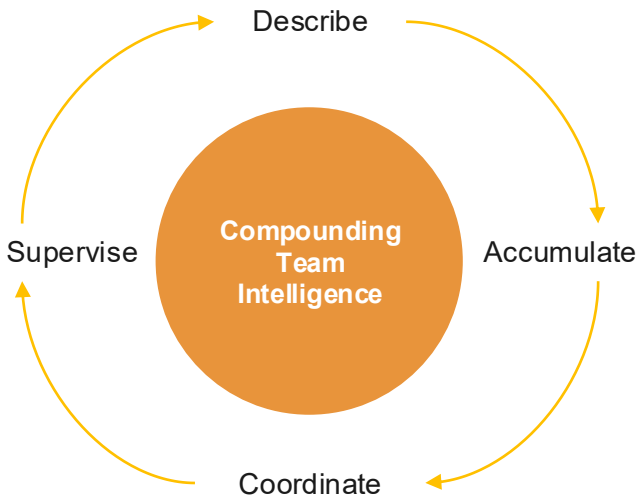

Positioning • Market • Competitive Landscape

Positioning

Not another chatbot. Not another agent framework.

Only platform where coordinated agent teams accumulate cross-platform work knowledge for autonomous output.

	yarnnn	Big-Players ChatGPT /Claude	Agent Startups Mem.ai, Genspark	Workspaces Notion, Slack
Coordinated agent teams	✓	✗	✗	✗
Cross-platform knowledge	✓ Evolving	✗	✗	Single
Cross-platform sync	✓	✗	✗	✗
Autonomous output	✓ Scheduled	✗	✓	✗
Improves with tenure	✓ Accumulation	Evolving	✗	✗
PM team coordination	✓	✗	✗	✗

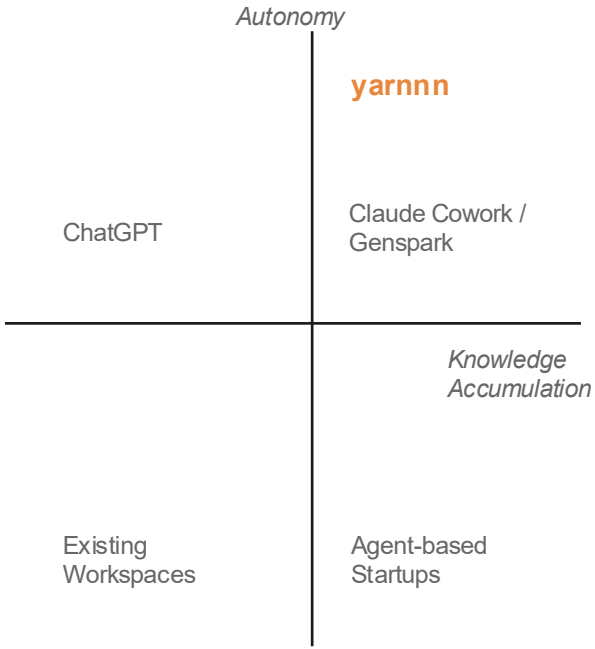

Market

Starting with solo consultants. Expanding from there.

Entry wedge: Solo consultants with recurring client obligations

TAM

$4.35B

AI productivity tools market (31% CAGR)

SAM

$1.14B

~5M solo consultants globally × $228/yr (Pro annual)

Entry SOM

$4.6M

20K consultants in first 3 years

Expansion path:

Now	**Solo consultants**	Recurring reports, clear pain, willing to pay
Next	**Founders & executives**	Investor updates, board prep, strategic briefs
Then	**Teams & ops leads**	Standups, project reports, cross-team synthesis
Scale	**All knowledge workers**	Any professional with recurring work obligations

Pricing Model

Two tiers. Work credits are the lever.

Free

$0 /mo

- 2 active tasks
- 20 work credits / month
- 150 chat messages / month
- 5 Slack / 10 Notion sources
- Daily sync

Pro

$19 /mo

- 10 active tasks
- 500 work credits / month
- Unlimited chat messages
- Unlimited sources
- Hourly sync

Enterprise

TBD /mo

- Everything in Pro
- Team workspaces
- Priority support
- Custom integrations
- Agent-to-agent protocol

Credits: task run = 3 credits, render/export = 1 credit. Multi-step pipelines cost 3 credits per step.

Unit economics: Target 70%+ gross margin at scale. Yearly plan: 2 months free (17% off). Pro $190/yr.

Product Roadmap

Coordinated Agent Teams (shipped) → Cross-Platform Intelligence → Agent Protocol

Each phase deepens the value chain. Shipped work compounds into the next layer of capability.

SHIPPED ✓

Coordinated Agent Teams

Work-first onboarding + project scaffolding

PM phase dispatch + quality gating

8 agent types (three-registry)

Output gateway: HTML compose, PDF, PPTX, XLSX

Feedback substrate: edit distillation + self-reflection

Workspace filesystem + cross-agent collaboration

136 Architecture Decision Records

NEXT 6-12mo.

Cross-Platform Intelligence

Projects: multi-agent assembled deliverables

Composer-driven assembly from contributions

Deeper platform coverage: 365, Google Drive

HTML-native compose as output pipeline

VISION +12mo.

Agent-to-Agent Protocol

MCP protocol interop. External agents consume yamnn context

Agent-to-agent coordination across systems

Platform play: own the work context graph

(Pre) Seed - The Ask

$500K at $5M post-money valuation.

Use of funds:

→ Tech Lead (memory/ML)

→ GTM Lead (creator community)

→ 12 months runway

Why now:

→ ClawdBot: 17,830 stars / 24hrs

→ No one owns the context layer

→ Architecture built, need team

Positioning:

→ Application layer still pre-mature

→ True Agent-native architecture

Priority 0: Team-Building

→ ML engineer for memory extraction

→ Community builder for creator GTM

→ Seed funds → team of 3

Betting on an agent-first future:

→ LLM readiness moving towards Agents

→ Thesis is the work economy will shift

Differentiation:

→ Coordinated teams, not solo agents

→ Cross-platform, LLM agnostic

→ 136 ADRs = architectural rigor

yarnnn

yarnnn.com · Kevin Kim · seulkim88@gmail.com

Kevin (Seul Ki) Kim

Founder & CEO, KVK Labs (yarnnn)

https://www.linkedin.com/in/kimseulki0727/

B.S. Industrial and Systems Engineering, Virginia Tech

+10 years designing CRM and GTM systems in FMCG.

CMO, CPO positions in Korea based startups, Ran e-commerce fashion brand for 5 years,

Full-stack application for yarnnn, (Next.js + FastAPI + Supabase), two platform integrations, coordinated agent team architecture with PM phase dispatch, documented across 136 Architecture Decision Records.

Targeting US market; execute GTM and team-building as priority next step.

Appendix: Value Realization Chain

What the User Experiences: Describe to Compound

Time	What happens	What the user sees
T+0s	User describes work, projects scaffold	"I've set up 3 projects for your client work"
T+30s	Platforms connect, sources mapped to projects	Dashboard: "Slack connected — channels mapped"
T+90s	PM dispatches first contributor runs	First deliverables appear per project
Day 2-7	Daily pulse cycles, agents accumulate	Deliverables arrive, improving each run
Week 2	PM coordination matures, phases sequence	Cross-phase handoffs visible in output
Month 1	Multiple agents running, feedback refining	Agents noticeably tailored to preferences
Month 3	Composer identifies cross-project opportunities	System understands work deeply

Describe → Scaffold → Connect → First Value → Accumulate → Coordinate → Compound ↻

What a Well-Funded Team Can't Shortcut

"Claude Code Online" — Output Gateway

Same capabilities model as Claude Code (SKILL.md + local tools + filesystem), but persistent, autonomous, and cloud-native. Two-filesystem architecture: capability filesystem (skills on Docker) + content filesystem (workspace per agent). Adding a new output format = adding a skill folder.

PM-Coordinated Phase Dispatch

PM owns the heartbeat. Contributors dispatched in sequence. Cross-phase context injection: researcher's findings → analyst's input. Quality-gated assembly. Not co-located agents — a coordinated team with a manager.

Feedback-to-Intelligence Loop

Three unified feedback rails: user edits distilled to persistent preferences, agent self-reflection after every run, PM reads contributor trajectories. "Improves with tenure" is structural, not aspirational.

Charter-Driven Deterministic Execution

Three charter files (PROJECT.md + TEAM.md + PROCESS.md) define execution deterministically. Cadence enforcement prevents runaway loops. ~$0.50/month per project at steady state. Predictable cost model.

Why this compounds into defensibility:

The PM coordination can't work without the feedback substrate. The feedback substrate can't work without the workspace filesystem. The workspace filesystem can't work without the perception pipeline populating it. A team starting today would need to discover these interdependencies through the same 136 decisions — or build a simpler system that can't compound.