yarnnn

Your autonomous AI agent — powered by accumulated work context

It connects to your tools. It learns your world. It works while you don't.

(Pre) Seed Round 2026

Kevin Kim · kvkthecreator@gmail.com

It looks like the big players will own everything.

Claude has coding agents, desktop, cowork.
ChatGPT has memory, Codex, Sora.
The market assumes they'll consume every layer.

But every platform cycle looks that way — until it doesn't.

Google didn't become Salesforce.
Facebook didn't become Shopify.
AWS didn't become Datadog.

The Honest Bet

History never repeats, but it does rhyme. Our thesis is that even with the AI tech revolution, the application layer is yet to come.

2008: Google	→	Looked like it would own everything on the web. It didn't become Salesforce.
2012: Facebook	→	Looked like it would own all of social commerce. It didn't become Shopify.
2015: AWS	→	Looked like it would own every cloud application. It didn't become Datadog.
2026: LLMs	→	Look like they'll own every layer. But code was the easy case. Work context is the hard one.

The application layer of the AI platform shift is yet to come, and we're taking the **bold bet** that it will be Agent-native, and the work economy will shift from human -> agents

Proof of Demand

ClawdBot proved the appetite is real.

17,830

GitHub stars in 24 hours

Fastest single-day growth in GitHub history.

> *"The first time I've felt like I'm living in the future since ChatGPT launched."*
>
> **— ClawdBot user**

But 95% of people can't use it.

High technical barrier, privacy, hardware setup

- ✗ Requires VPS provisioning (only 5% can do this)
- ✗ 200+ security leaks detected (GitGuardian)
- ✗ No project isolation or recurring workflows

Product Architecture Solution / Moat

Product: meet yarnnn’s agent: Thinking Partner TP

Your autonomous AI agent. Live at yarnnn.com.

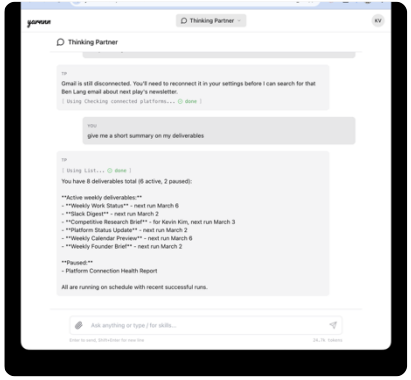

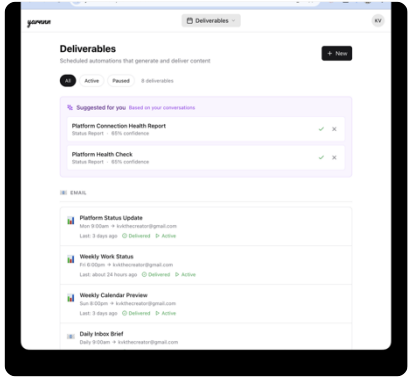

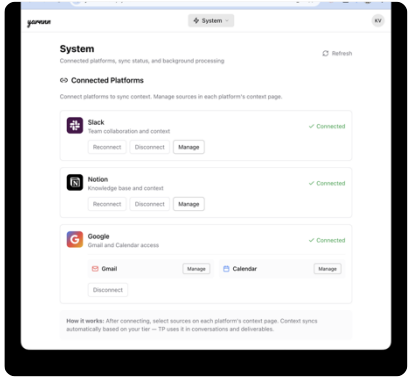

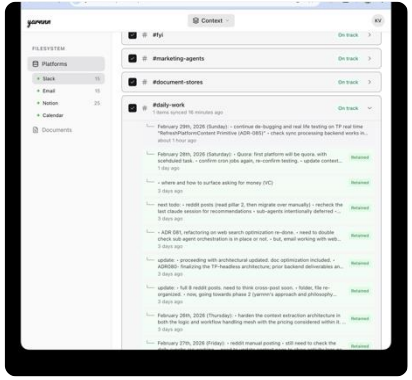

01 Chat Mode

Ask anything. TP already has the context. Built like Claude Code with primitive tools, webSearch

02 Headless Mode

Autonomous deliverables on schedule. Backend Processors, Cron jobs to make autonomy native

03 4 Integrations

Slack, Gmail, Notion, Calendar. With more in product roadmap

04 Context Engine

Gets smarter every sync cycle. Recursive context management, memory handling

Agent Architecture and Orchestration

Built for Agent-native, autonomous workflows

Designed so context compounds over time, creating recursive feedback between outputs and agent architecture

Platform Sync
Slack • Gmail
Notion • Calendar

→

Content Layer
Unified store
Retention + provenance

→

Memory
Working + user memory
Implicit extraction

→

Agent Engine
5 execution modes
Chat + headless

→

Deliverables
Autonomous
scheduled output

Why this can’t be bolted on:

Big Players, LLM Providers
No persistent cross-platform context.
Optimized for breadth, not individual depth.

Agent Startups
Architectural gap.
Task execution focused.
No accumulated context layer with feedback loop

Workspace AI
Single-platform. Can’t synthesize across tools.
No autonomous deliverables.

Every layer purpose-built for context accumulation, true autonomy + intelligence

Solution: Compounding Stickiness

Value Compounds with Every Execution Cycle.

Connect your platforms. TP learns your world. Deliverables run on their own.

DAY 1

Connect + Context

Link Slack, Gmail, Notion, Calendar. First sync in minutes. TP already knows your work.

WHAT YOU GET:

- Ask TP anything — shared context is key
- Pre-meeting brief from Slack + Gmail
- Project status from channel activity

DAY 30

Autonomous Output

Scheduled deliverables run without prompting. Each one learns from your edits and feedback.

WHAT RUNS ON ITS OWN:

- Weekly client updates from project channels
- Monday meeting prep with attendee context
- Monthly board report from all platforms

DAY 90

Irreplaceable Intelligence

90 days of accumulated context. TP knows your clients, your projects, your style. No generic tool can replicate this.

WHAT ONLY YOU HAVE:

- Deal alerts when pricing surfaces in Slack
- Cross-platform research digests
- Switching cost: context is irreplaceable

Solution: Moat

Switching Costs compound as Context Accumulates

Every interaction makes yarnnn harder to replace and more valuable to keep.

Why This Creates a Moat:

1. Accumulation: Every sync cycle deepens cross-platform understanding
2. Recursion: Your edits teach the system — output quality compounds
3. Switching costs: 90 days of context is irreplaceable by any competitor
4. Lock-in: Deliverables depend on accumulated context — not transferable

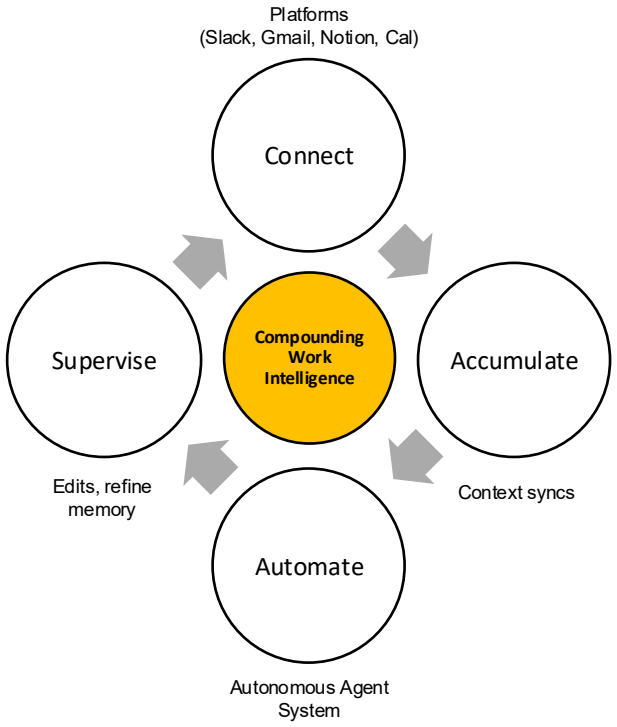

Proactive Intelligence

Your AI works while you don't.

Two-phase execution: cheap review decides if expensive generation is warranted.

Phase 1: Review
Haiku scans accumulated context
Cost: ~$0.002 per cycle
Runs every scheduled interval

→

Decision Gate
Is there something worth reporting?
Most cycles: nothing new
Only triggers generation when signal detected

→

Phase 2: Generate
Sonnet crafts personalized output
Cost: ~$0.05 per generation
Only fires when review finds signal

25x cost difference means we can monitor continuously without burning budget.

Positioning • Market • Competitive Landscape

Positioning

Not another chatbot. Not another agent framework.

The only agent that accumulates cross-platform work context and uses it for autonomous output.

	yarnnn	**Big-Players** ChatGPT/Claude	**Agent Startups** Mem.ai, Genspark	**Workspaces** Notion, Slack
Cross-platform context accumulation	✓ Evolving	✗	✗	Single
Cross-platform sync	✓	✗	✗	✗
Autonomous output	✓ Scheduled	✗	✓	✗
Improves with tenure	✓ Accumulation	Evolving	✗	✗
Context-aware agent	✓	Evolving	✗	Basic

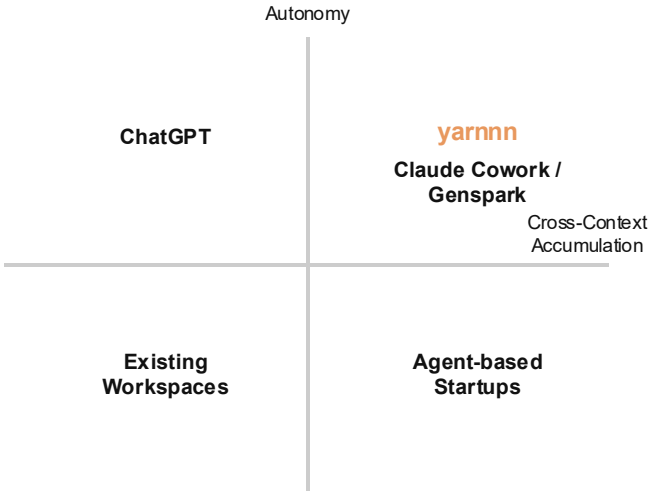

Market

Starting with solo consultants. Expanding from there.

Entry wedge: Solo consultants with recurring client obligations

TAM
$4.35B
AI productivity tools market (31% CAGR)

SAM
$1.14B
~5M solo consultants globally × $228/yr (Pro annual)

Entry SOM
$11.4M
50K consultants (1% SAM) in first 3 years

Expansion path:

Now	**Solo consultants**	Recurring reports, clear pain, willing to pay
Next	**Founders & executives**	Investor updates, board prep, strategic briefs
Then	**Teams & ops leads**	Standups, project reports, cross-team synthesis
Scale	**All knowledge workers**	Any professional with recurring work obligations

Pricing Model

Three tiers. Sync frequency is the lever.

Free

$0 /mo

- 2 platforms connected
- 1 source per platform
- Sync 2x/day
- 20 TP conversations/mo
- 3 active deliverables

Starter

$9 /mo

- 4 platforms connected
- 5 sources per platform
- Sync every 6 hours
- 100 TP conversations/mo
- 10 active deliverables

Pro

$19 /mo

- 4 platforms connected
- 20+ sources per platform
- Hourly sync
- Unlimited TP conversations
- Unlimited deliverables

Unit economics: Target 70%+ gross margin at scale. Batched sync and cached context reduce per-user API cost as context accumulates.

Yearly plans: 2 months free (17% off). Starter $90/yr, Pro $190/yr.

Product Roadmap

Context layer first. Autonomy second. Protocol third.

Each phase expands the moat. More context surfaces unlock deeper autonomy and broader interoperability.

NEAR-TERM <6mo.

Expand the Context Surface

Linear, GitHub, MS Office, Google Drive
Every new integration deepens the accumulation moat
Cross-platform context is what makes deliverables intelligent

MEDIUM-TERM 12-18mo.

Long-term Autonomy Agent

Move from proactive review to full autonomous execution
Agent drafts, schedules, and acts on your behalf with accumulated context
Trust earned over time: the more it knows you, the more you delegate

LONG-TERM +12mo.

Agent-to-Agent Protocol

YARNNN becomes the context layer other agents talk to
Interop with external AI agents via standardized protocol
Platform play: own the work context graph, not just one agent

(Pre) Seed - The Ask

$500K–$1M seed round at $5–10M valuation.

Use of funds:

→ Tech Lead (memory/ML)

→ GTM Lead (creator community)

→ 12–18 months runway

Why now:

→ ClawdBot: 17,830 stars / 24hrs

→ No one owns the context layer

→ Architecture built, need team

Positioning:

→ Application layer sill pre-mature

→ True Agent-native architecture

Priority 0: Team-Building

→ ML engineer for memory extraction

→ Community builder for creator GTM

→ Seed funds → team of 3

Betting on an agent-first future:

→ LLM readiness moving towards Agents

→ Thesis is the work economy will shift

Differentiation:

→ Cross-platform, LLM agnostic

→ Big players can’t do cross data accumulations

yarnnn

yarnnn.com · Kevin Kim · kvkthecreator@gmail.com

Appendix

Kevin Kim

Founder & CEO, KVK Labs (yarnnn)

https://www.linkedin.com/in/kimseulki0727/

B.S. Industrial and Systems Engineering, Virginia Tech

+10 years designing CRM and GTM systems in FMCG.

CMO, CPO positions in Korea based startups, Ran e-commerce fashion brand for 5 years,

Full-stack application for yarnnn, (Next.js + FastAPI + Supabase), four platform integrations, unified agent architecture documented across 90+ Architecture Decision Records.

Targeting US market; execute GTM and team-building as priority next step.

kvkthecreator@gmail.com · yarnnn.com