yarnnn

Autonomous agents that know your work.

They connect to your platforms. They accumulate knowledge. They work while you don't.

(Pre) Seed Round 2026

Kevin Kim · kvkthecreator@gmail.com

It looks like the big players will own everything.

Claude has coding agents, desktop, cowork.
ChatGPT has memory, Codex, Sora.
The market assumes they'll consume every layer.

Every platform cycle looks that way — until it doesn't.

Google didn't become Salesforce.
Facebook didn't become Shopify.
AWS didn't become Datadog.

The Honest Bet

History never repeats, but it does rhyme. Our thesis is that even with the AI tech revolution, the application layer is yet to come.

2008: Google	→	Looked like it would own everything on the web. It didn't become Salesforce.
2012: Facebook	→	Looked like it would own all of social commerce. It didn't become Shopify.
2015: AWS	→	Looked like it would own every cloud application. It didn't become Datadog.
2026: LLMs	→	Look like they'll own every layer. But code was the easy case. Work context is the hard one.

The application layer of the AI platform shift is yet to come, and we're taking the **bold bet** that it will be Agent-native, and the work economy will shift from human -> agents

Proof of Demand

ClawdBot proved the appetite is real.

17,830

GitHub stars in 24 hours

Fastest single-day growth in GitHub history.

> *"The first time I've felt like I'm living in the future since ChatGPT launched."*
>
> **— ClawdBot user**

But 95% of people can't use it.

High technical barrier, privacy, hardware setup

- ✗ Requires VPS provisioning (only 5% can do this)
- ✗ 200+ security leaks detected (GitGuardian)
- ✗ No project isolation or recurring workflows

yarnnn

Product
Architecture
Solution / Moat

Product: Meet yarnnn's Orchestrator

Your autonomous agent network. Live at yarnnn.com.

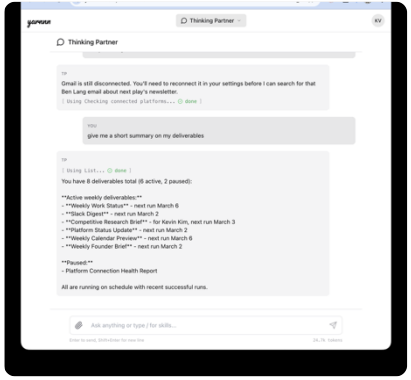

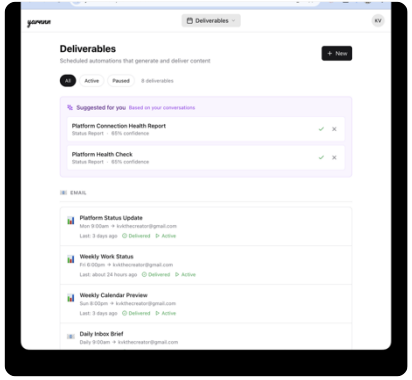

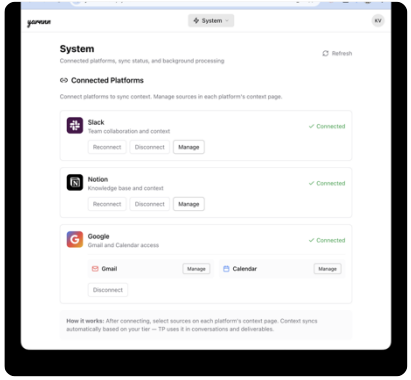

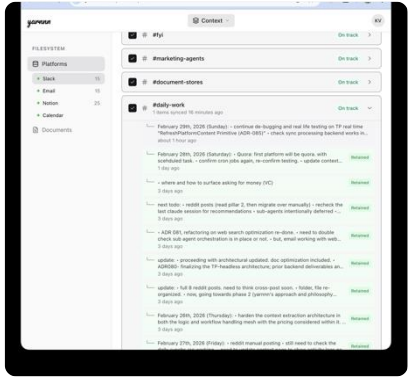

01 Chat Mode

Ask anything. It already knows your work. Built like Claude Code with capabilities and webSearch

02 Headless Mode

Autonomous agents on schedule. Persistent agents with inspectable workspaces

03 4 Integrations

Slack, Gmail, Notion, Calendar. With more in product roadmap

04 Context Engine

Gets smarter every sync cycle. Knowledge accumulation, workspace memory

Agent Architecture and Orchestration

Built for Agent-native, autonomous workflows

Every layer purpose-built so knowledge compounds — creating recursive feedback between agent output and accumulated intelligence

Perception Pipeline
Slack • Gmail
Notion • Calendar

→

Context Mgmt.
Shared content management for multi agents
Retention + provenance

→

Memory
Working + user memory
Implicit extraction

→

Agent Orchestration
5 execution modes
Chat + headless (cron-jobs, sync processors, etc.)

→

Agents
Persistent autonomous output

Why this can’t be bolted on:

Big Players, LLM Providers
Cross-platform knowledge not possible.
Optimized for horizontal, not individual depth.

Agent Startups (mem.ai, openclaw)
Architectural gap. (hardware or agent framework)
No accumulated knowledge layer with feedback loop

Workspace AI (Notion AI, Slack)
Single-platform. Can’t synthesize across tools.
No persistent agents.

Product

Your Agent Network

Platform specialists accumulate knowledge. Synthesis agents produce work you'd never get from a single tool.

RECAP AGENT

Platform Digest

What happened while you weren't looking. All platforms synthesized into one brief.

EXAMPLE:

Monday morning: here's everything from #engineering, your inbox, and 3 Notion updates. Gets better every week.

MEETING PREP AGENT

Calendar-Driven Briefs

Full attendee context from Slack + Gmail + Notion. Delivered before every meeting.

EXAMPLE:

You meet Sarah at 2pm. Here's her last 5 Slack threads, 2 emails, and the Notion doc she edited yesterday.

STATUS + WATCH AGENTS

Synthesis + Proactive Intel

Cross-platform status updates for any audience. Proactive alerts when pricing or competitors surface across channels.

EXAMPLES:

Weekly client update from Slack + Gmail + Notion. No copy-paste. Competitor mention detected across 3 channels. Brief generated.

Solution: Moat

Knowledge Compounds. Switching Costs Follow.

Every agent run makes the next one better.
Every edit teaches the system.
Every sync deepens the knowledge base.

Three moat mechanics:

1. Accumulation: Every sync cycle deepens cross-platform understanding. Agents don't start from zero.
2. Recursion: Your edits and feedback train agent output quality. The 5th run needs fewer corrections than the 1st.
3. Irreversibility: 90 days of accumulated knowledge can't be exported or replicated. Switching means starting over.

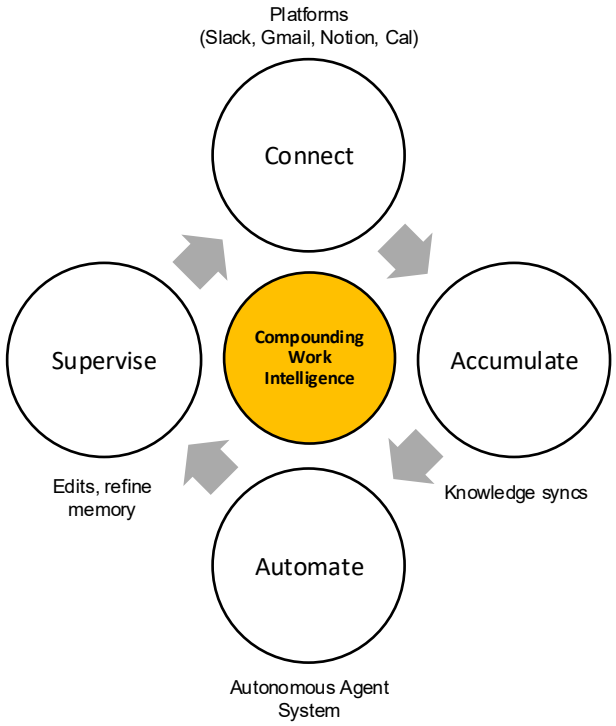

Positioning • Market • Competitive Landscape

Positioning

Not another chatbot. Not another agent framework.

The only platform where persistent agents accumulate cross-platform work knowledge for autonomous output.

	yarnnn	**Big-Players** ChatGPT/Claude	**Agent Startups** Mem.ai, Genspark	**Workspaces** Notion, Slack
Cross-platform knowledge accumulation	✓ Evolving	✗	✗	Single
Cross-platform sync	✓	✗	✗	✗
Autonomous output	✓ Scheduled	✗	✓	✗
Improves with tenure	✓ Accumulation	Evolving	✗	✗
Context-aware agent	✓	Evolving	✗	Basic

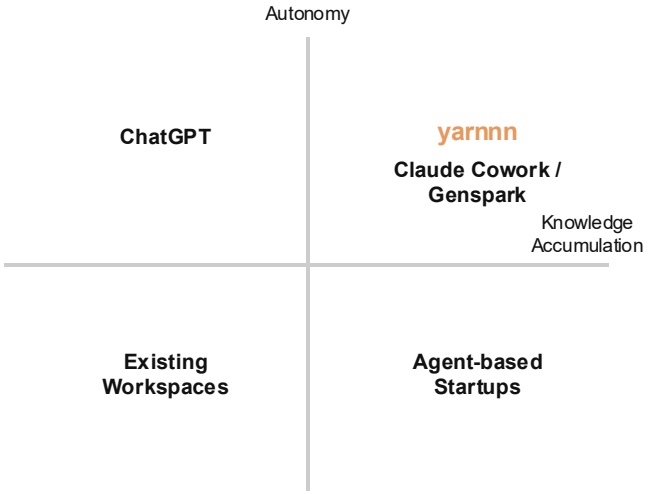

Market

Starting with solo consultants. Expanding from there.

Entry wedge: Solo consultants with recurring client obligations

TAM
$4.35B
AI productivity tools market (31% CAGR)

SAM
$1.14B
~5M solo consultants globally × $228/yr (Pro annual)

Entry SOM
$4.6M
20K consultants in first 3 years

Expansion path:

Now	**Solo consultants**	Recurring reports, clear pain, willing to pay
Next	**Founders & executives**	Investor updates, board prep, strategic briefs
Then	**Teams & ops leads**	Standups, project reports, cross-team synthesis
Scale	**All knowledge workers**	Any professional with recurring work obligations

Pricing Model

Two tiers. Knowledge depth is the lever.

Free

$0 /mo

- 2 platforms connected
- 5 sources per platform
- Sync 2x/day
- 50 orchestrator messages/mo
- 2 active agents

Pro

$19 /mo

- 4 platforms connected
- Unlimited sources
- Hourly sync
- Unlimited orchestrator messages
- Unlimited agents

Enterprise

TBD /mo

- Everything in Pro
- Team workspaces
- Priority support
- Custom integrations
- Agent-to-agent protocol

Unit economics: Target 70%+ gross margin at scale. Batched sync and cached knowledge reduce per-user API cost as context accumulates.

Yearly plan: 2 months free (17% off). Pro $190/yr.

Product Roadmap

Context layer first. Autonomy second. Protocol third.

Each phase expands the moat. More context surfaces unlock deeper autonomy and broader interoperability.

NEAR-TERM <6mo.

Expand the Context Surface

Linear, GitHub, MS Office, Google Drive
Every new integration deepens the accumulation moat
Cross-platform knowledge is what makes agents intelligent

MEDIUM-TERM 12-18mo.

Long-term Autonomy Agent

Move from proactive review to full autonomous execution
Agent drafts, schedules, and acts on your behalf with accumulated context
Trust earned over time: the more it knows you, the more you delegate

LONG-TERM +12mo.

Agent-to-Agent Protocol

YARNNN becomes the context layer other agents talk to
Interop with external AI agents via standardized protocol
Platform play: own the work context graph, not just one agent

(Pre) Seed - The Ask

$500K at $5M post-money valuation.

Use of funds:

→ Tech Lead (memory/ML)

→ GTM Lead (creator community)

→ 12 months runway

Why now:

→ ClawdBot: 17,830 stars / 24hrs

→ No one owns the context layer

→ Architecture built, need team

Positioning:

→ Application layer sill pre-mature

→ True Agent-native architecture

Priority 0: Team-Building

→ ML engineer for memory extraction

→ Community builder for creator GTM

→ Seed funds → team of 3

Betting on an agent-first future:

→ LLM readiness moving towards Agents

→ Thesis is the work economy will shift

Differentiation:

→ Cross-platform, LLM agnostic

→ Big players can’t do cross data accumulations

yarnnn

yarnnn.com · Kevin Kim · kvkthecreator@gmail.com

Appendix

Kevin Kim

Founder & CEO, KVK Labs (yarnnn)

https://www.linkedin.com/in/kimseulki0727/

B.S. Industrial and Systems Engineering, Virginia Tech

+10 years designing CRM and GTM systems in FMCG.

CMO, CPO positions in Korea based startups, Ran e-commerce fashion brand for 5 years,

Full-stack application for yarnnn, (Next.js + FastAPI + Supabase), four platform integrations, unified agent architecture documented across 115+ Architecture Decision Records.

Targeting US market; execute GTM and team-building as priority next step.

kvkthecreator@gmail.com · yarnnn.com

Product

What YARNNN Replaces

Three real workflows. Before: manual, fragmented, time-consuming. After: autonomous, cross-platform, improving.

Before: Monday Catchup

Open Slack. Scroll 14 channels. Open Gmail. Scan 40 emails. 45 minutes gone.

AFTER: Recap Agent
One brief. All platforms. Waiting at 8am. 0 minutes of your time.

Before: Client Status Update

Copy from Slack. Paste into doc. Check Gmail. Rewrite for client tone. 2 hours, every week.

AFTER: Status Agent
Pulls from all platforms. 12 weeks of accumulated context. Draft arrives Friday 5pm.

Before: Meeting Prep

Search Slack. Search Gmail. Check LinkedIn. Skim Notion. Scramble 10 minutes before the call.

AFTER: Meeting Prep Agent
Brief auto-generated from calendar. Full context. Delivered 30 min before. Every time.

Every agent gets better with tenure. The 12th weekly update needs zero corrections.